

Australian dataset classification

March 23, 2023

```
[1]: import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LogisticRegression
      from sklearn.tree import DecisionTreeClassifier
```

```
[2]: from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
      from sklearn.neural_network import MLPClassifier
      from sklearn.metrics import accuracy_score, f1_score, precision_score, \
      ↪ recall_score
```

```
[3]: import warnings
      warnings.filterwarnings("ignore")
```

```
[4]: data = pd.read_csv('australian.csv')
```

```
[5]: data.head(2)
```

```
[5]:  @inputs A1    A2    A3  A4  A5  A6    A7  A8  A9  A10  A11  A12  A13  A14  \
0          1  2208  1146   2   4   4  1585   0   0   0    1    2   100  1213
1          0  2267    7   2   8   4   165   0   0   0    0    2   160    1

      @output Class
0              0
1              0
```

```
[6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 690 entries, 0 to 689
Data columns (total 15 columns):
 #   Column          Non-Null Count  Dtype
---  -
0   @inputs A1      690 non-null   int64
1   A2              690 non-null   int64
2   A3              690 non-null   int64
3   A4              690 non-null   int64
4   A5              690 non-null   int64
5   A6              690 non-null   int64
6   A7              690 non-null   int64
```

```

7   A8          690 non-null    int64
8   A9          690 non-null    int64
9   A10         690 non-null    int64
10  A11         690 non-null    int64
11  A12         690 non-null    int64
12  A13         690 non-null    int64
13  A14         690 non-null    int64
14  @output Class 690 non-null    int64
dtypes: int64(15)
memory usage: 81.0 KB

```

```
[7]: data = data.rename(columns={"@inputs A1": "A1", "@output Class": "Class"})
```

```
[8]: data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 690 entries, 0 to 689
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   A1          690 non-null    int64
1   A2          690 non-null    int64
2   A3          690 non-null    int64
3   A4          690 non-null    int64
4   A5          690 non-null    int64
5   A6          690 non-null    int64
6   A7          690 non-null    int64
7   A8          690 non-null    int64
8   A9          690 non-null    int64
9   A10         690 non-null    int64
10  A11         690 non-null    int64
11  A12         690 non-null    int64
12  A13         690 non-null    int64
13  A14         690 non-null    int64
14  Class       690 non-null    int64
dtypes: int64(15)
memory usage: 81.0 KB

```

```

[9]: # Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(data.iloc[:, :-1], data.
    ↪iloc[:, -1], test_size=0.2, random_state=42)

```

0.0.1 Logistic Regression

```

[10]: # Train and evaluate a Logistic Regression model
lr_model = LogisticRegression(random_state=42)

```

```
[11]: lr_model.fit(X_train, y_train)
lr_preds = lr_model.predict(X_test)
```

```
[12]: lr_acc = accuracy_score(y_test, lr_preds)
lr_prec = precision_score(y_test, lr_preds)
lr_rec = recall_score(y_test, lr_preds)
lr_f1 = f1_score(y_test, lr_preds)
```

```
[13]: print("Logistic Regression Accuracy:", lr_acc)
print("Logistic Regression Precision:", lr_prec)
print("Logistic Regression Recall:", lr_rec)
print("Logistic Regression F1 Score:", lr_f1)
```

Logistic Regression Accuracy: 0.7608695652173914
Logistic Regression Precision: 0.6666666666666666
Logistic Regression Recall: 0.7058823529411765
Logistic Regression F1 Score: 0.6857142857142857

```
[14]: import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix
import numpy as np
import itertools
```


```
[15]: # Define class labels
classes = ['Class A', 'Class B', 'Class C']
```

```
[16]: # Compute confusion matrix
cm = confusion_matrix(y_test, lr_preds)
```

```
[17]: # Plot confusion matrix
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion matrix')
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)
plt.xlabel('Predicted value')
plt.ylabel('True value')

# Add text to each cell
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], 'd'),
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.show()
```



0.0.2 Decision Tree

```
[18]: # Train and evaluate a Decision Tree model
dt_model = DecisionTreeClassifier(random_state=42)
dt_model.fit(X_train, y_train)
```

```
[18]: DecisionTreeClassifier(random_state=42)
```

```
[19]: dt_preds = dt_model.predict(X_test)
dt_acc = accuracy_score(y_test, dt_preds)
dt_prec = precision_score(y_test, dt_preds)
dt_rec = recall_score(y_test, dt_preds)
dt_f1 = f1_score(y_test, dt_preds)
```

```
[20]: print("Results of Decision Tree:")
print("Decision Tree Accuracy:", dt_acc)
print("Decision Tree Precision:", dt_prec)
print("Decision Tree Recall:", dt_rec)
print("Decision Tree F1 Score:", dt_f1)
```

```
Results of Decision Tree:
Decision Tree Accuracy: 0.8333333333333334
Decision Tree Precision: 0.7692307692307693
Decision Tree Recall: 0.7843137254901961
```

Decision Tree F1 Score: 0.7766990291262137


```
[21]: import matplotlib.pyplot as plt
      from sklearn.metrics import roc_curve, auc

[22]: # Calculate predicted probabilities for positive class
      dt_probs = dt_model.predict_proba(X_test)[:, 1]

[23]: # Calculate FPR, TPR, and thresholds
      fpr, tpr, thresholds = roc_curve(y_test, dt_probs)

[24]: # Calculate AUC
      auc_dt = auc(fpr, tpr)

[25]: # Plot ROC curve
      plt.plot(fpr, tpr, label='Decision Tree (AUC = {:.2f})'.format(auc_dt))
      plt.plot([0, 1], [0, 1], linestyle='--', color='gray', label='Random Guess')
      plt.xlabel('False Positive Rate')
      plt.ylabel('True Positive Rate')
      plt.title('ROC Curve')
      plt.legend()
      plt.show()
```



0.0.3 Gradient Boosting

```
[26]: gb_model = GradientBoostingClassifier(random_state=42)
      gb_model.fit(X_train, y_train)
```

```
[26]: GradientBoostingClassifier(random_state=42)
```

```
[27]: gb_preds = gb_model.predict(X_test)
```

```
[28]: gb_acc = accuracy_score(y_test, gb_preds)
      gb_prec = precision_score(y_test, gb_preds)
      gb_rec = recall_score(y_test, gb_preds)
      gb_f1 = f1_score(y_test, gb_preds)
```

```
[29]: print("Results of Gradient Boosting")
      print("Gradient Boosting Accuracy:", gb_acc)
      print("Gradient Boosting Precision:", gb_prec)
      print("Gradient Boosting Recall:", gb_rec)
      print("Gradient Boosting F1 Score:", gb_f1)
```

Results of Gradient Boosting

Gradient Boosting Accuracy: 0.8695652173913043

Gradient Boosting Precision: 0.8367346938775511

Gradient Boosting Recall: 0.803921568627451

Gradient Boosting F1 Score: 0.8200000000000001

0.0.4 Artificial Neural Network

```
[30]: from sklearn.preprocessing import StandardScaler
      from keras.models import Sequential
      from keras.layers import Dense
```

```
[31]: # Feature Scaling
      sc = StandardScaler()
      X_train = sc.fit_transform(X_train)
      X_test = sc.transform(X_test)
```

```
[32]: # Define the ANN model
      model = Sequential()
      model.add(Dense(units=16, activation='relu', input_dim=X_train.shape[1]))
      model.add(Dense(units=8, activation='relu'))
      model.add(Dense(units=1, activation='sigmoid'))
```

```
[33]: # Compile the model
      model.compile(optimizer='adam', loss='binary_crossentropy',
                    metrics=['accuracy'])
```

```
[34]: # Train the model
      model.fit(X_train, y_train, batch_size=32, epochs=50)
```

```
Epoch 1/50
18/18 [=====] - 1s 2ms/step - loss: 0.7125 - accuracy:
0.5344
Epoch 2/50
18/18 [=====] - 0s 2ms/step - loss: 0.6512 - accuracy:
0.6757
Epoch 3/50
18/18 [=====] - 0s 2ms/step - loss: 0.6018 - accuracy:
0.7174
Epoch 4/50
18/18 [=====] - 0s 2ms/step - loss: 0.5588 - accuracy:
0.7518
Epoch 5/50
18/18 [=====] - 0s 2ms/step - loss: 0.5200 - accuracy:
0.7826
Epoch 6/50
18/18 [=====] - 0s 2ms/step - loss: 0.4844 - accuracy:
0.8043
Epoch 7/50
18/18 [=====] - 0s 2ms/step - loss: 0.4551 - accuracy:
0.8207
Epoch 8/50
18/18 [=====] - 0s 2ms/step - loss: 0.4280 - accuracy:
0.8406
Epoch 9/50
18/18 [=====] - 0s 1ms/step - loss: 0.4053 - accuracy:
0.8496
Epoch 10/50
18/18 [=====] - 0s 1ms/step - loss: 0.3858 - accuracy:
0.8605
Epoch 11/50
18/18 [=====] - 0s 1ms/step - loss: 0.3706 - accuracy:
0.8641
Epoch 12/50
18/18 [=====] - 0s 2ms/step - loss: 0.3580 - accuracy:
0.8678
Epoch 13/50
18/18 [=====] - 0s 2ms/step - loss: 0.3470 - accuracy:
0.8750
Epoch 14/50
18/18 [=====] - 0s 2ms/step - loss: 0.3378 - accuracy:
0.8822
Epoch 15/50
18/18 [=====] - 0s 2ms/step - loss: 0.3306 - accuracy:
0.8822
```

Epoch 16/50
18/18 [=====] - 0s 2ms/step - loss: 0.3237 - accuracy:
0.8859
Epoch 17/50
18/18 [=====] - 0s 2ms/step - loss: 0.3182 - accuracy:
0.8895
Epoch 18/50
18/18 [=====] - 0s 2ms/step - loss: 0.3133 - accuracy:
0.8931
Epoch 19/50
18/18 [=====] - 0s 1ms/step - loss: 0.3088 - accuracy:
0.8913
Epoch 20/50
18/18 [=====] - 0s 1ms/step - loss: 0.3052 - accuracy:
0.8913
Epoch 21/50
18/18 [=====] - 0s 1ms/step - loss: 0.3014 - accuracy:
0.8877
Epoch 22/50
18/18 [=====] - 0s 1ms/step - loss: 0.2985 - accuracy:
0.8913
Epoch 23/50
18/18 [=====] - 0s 2ms/step - loss: 0.2959 - accuracy:
0.8913
Epoch 24/50
18/18 [=====] - 0s 2ms/step - loss: 0.2927 - accuracy:
0.8967
Epoch 25/50
18/18 [=====] - 0s 2ms/step - loss: 0.2897 - accuracy:
0.8967
Epoch 26/50
18/18 [=====] - 0s 2ms/step - loss: 0.2872 - accuracy:
0.8967
Epoch 27/50
18/18 [=====] - 0s 2ms/step - loss: 0.2849 - accuracy:
0.8967
Epoch 28/50
18/18 [=====] - 0s 2ms/step - loss: 0.2827 - accuracy:
0.8967
Epoch 29/50
18/18 [=====] - 0s 1ms/step - loss: 0.2801 - accuracy:
0.9004
Epoch 30/50
18/18 [=====] - 0s 2ms/step - loss: 0.2785 - accuracy:
0.9004
Epoch 31/50
18/18 [=====] - 0s 2ms/step - loss: 0.2768 - accuracy:
0.9022

Epoch 32/50
18/18 [=====] - 0s 1ms/step - loss: 0.2753 - accuracy:
0.9022
Epoch 33/50
18/18 [=====] - 0s 2ms/step - loss: 0.2725 - accuracy:
0.9040
Epoch 34/50
18/18 [=====] - 0s 2ms/step - loss: 0.2715 - accuracy:
0.9022
Epoch 35/50
18/18 [=====] - 0s 2ms/step - loss: 0.2697 - accuracy:
0.9022
Epoch 36/50
18/18 [=====] - 0s 2ms/step - loss: 0.2680 - accuracy:
0.9040
Epoch 37/50
18/18 [=====] - 0s 2ms/step - loss: 0.2662 - accuracy:
0.9040
Epoch 38/50
18/18 [=====] - 0s 2ms/step - loss: 0.2651 - accuracy:
0.9040
Epoch 39/50
18/18 [=====] - 0s 1ms/step - loss: 0.2634 - accuracy:
0.9058
Epoch 40/50
18/18 [=====] - 0s 1ms/step - loss: 0.2620 - accuracy:
0.9040
Epoch 41/50
18/18 [=====] - 0s 2ms/step - loss: 0.2612 - accuracy:
0.9040
Epoch 42/50
18/18 [=====] - 0s 1ms/step - loss: 0.2604 - accuracy:
0.9058
Epoch 43/50
18/18 [=====] - 0s 1ms/step - loss: 0.2586 - accuracy:
0.9076
Epoch 44/50
18/18 [=====] - 0s 1ms/step - loss: 0.2583 - accuracy:
0.9076
Epoch 45/50
18/18 [=====] - 0s 1ms/step - loss: 0.2567 - accuracy:
0.9094
Epoch 46/50
18/18 [=====] - 0s 1ms/step - loss: 0.2548 - accuracy:
0.9094
Epoch 47/50
18/18 [=====] - 0s 1ms/step - loss: 0.2535 - accuracy:
0.9094

```
Epoch 48/50
18/18 [=====] - 0s 1ms/step - loss: 0.2522 - accuracy:
0.9076
Epoch 49/50
18/18 [=====] - 0s 2ms/step - loss: 0.2515 - accuracy:
0.9149
Epoch 50/50
18/18 [=====] - 0s 1ms/step - loss: 0.2505 - accuracy:
0.9094
```

```
[34]: <keras.callbacks.History at 0x1afa656fa00>
```

```
[35]: # Evaluate the model
y_pred = model.predict(X_test)
y_pred = (y_pred > 0.5).astype(int)
```

```
5/5 [=====] - 0s 1ms/step
```

```
[36]: accuracy = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
```

```
[37]: print("Results of Artificial Neural Network:")
print("ANN Accuracy:", accuracy)
print("ANN Precision:", prec)
print("ANN Recall:", rec)
print("ANN F1 Score:", f1)
```

```
Results of Artificial Neural Network:
ANN Accuracy: 0.8840579710144928
ANN Precision: 0.8888888888888888
ANN Recall: 0.7843137254901961
ANN F1 Score: 0.8333333333333334
```