

heart classification

March 24, 2023

```
[1]: import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LogisticRegression
      from sklearn.tree import DecisionTreeClassifier
```

```
[2]: from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
      from sklearn.neural_network import MLPClassifier
      from sklearn.metrics import accuracy_score, f1_score, precision_score, \
      ↪ recall_score
```

```
[3]: import warnings
      warnings.filterwarnings("ignore")
```

```
[4]: data = pd.read_csv('heart.csv')
```

```
[5]: data.head(2)
```

```
[5]: @inputs Age Sex ChestPainType RestBloodPressure SerumCholestoral \
0      70 1 4 130 322
1      67 0 3 115 564

      FastingBloodSugar ResElectrocardiographic MaxHeartRate ExerciseInduced \
0      0 2 109 0
1      0 2 160 0

      Oldpeak Slope MajorVessels Thal Class
0      24 2 3 3 2
1      16 2 0 7 1
```

```
[6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 270 entries, 0 to 269
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   @inputs Age           270 non-null   int64
1   Sex                   270 non-null   int64
2   ChestPainType         270 non-null   int64
```

```

3   RestBloodPressure      270 non-null   int64
4   SerumCholestoral       270 non-null   int64
5   FastingBloodSugar      270 non-null   int64
6   ResElectrocardiographic 270 non-null   int64
7   MaxHeartRate           270 non-null   int64
8   ExerciseInduced        270 non-null   int64
9   Oldpeak                270 non-null   int64
10  Slope                  270 non-null   int64
11  MajorVessels           270 non-null   int64
12  Thal                   270 non-null   int64
13  Class                  270 non-null   int64

```

dtypes: int64(14)

memory usage: 29.7 KB

```
[7]: data = data.rename(columns={"@inputs Age": "Age"})
```

```
[8]: data1 = data
```

```
[9]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 270 entries, 0 to 269

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	Age	270 non-null	int64
1	Sex	270 non-null	int64
2	ChestPainType	270 non-null	int64
3	RestBloodPressure	270 non-null	int64
4	SerumCholestoral	270 non-null	int64
5	FastingBloodSugar	270 non-null	int64
6	ResElectrocardiographic	270 non-null	int64
7	MaxHeartRate	270 non-null	int64
8	ExerciseInduced	270 non-null	int64
9	Oldpeak	270 non-null	int64
10	Slope	270 non-null	int64
11	MajorVessels	270 non-null	int64
12	Thal	270 non-null	int64
13	Class	270 non-null	int64

dtypes: int64(14)

memory usage: 29.7 KB

```
[19]: # Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(data.iloc[:, :-1], data.
    ↪iloc[:, -1], test_size=0.2, random_state=42)
```

0.0.1 Logistic Regression

```
[20]: # Train and evaluate a Logistic Regression model
lr_model = LogisticRegression(random_state=42)
```

```
[21]: lr_model.fit(X_train, y_train)
lr_preds = lr_model.predict(X_test)
```

```
[22]: lr_acc = accuracy_score(y_test, lr_preds)
lr_prec = precision_score(y_test, lr_preds)
lr_rec = recall_score(y_test, lr_preds)
lr_f1 = f1_score(y_test, lr_preds)
```

```
[23]: print("Logistic Regression Accuracy:", lr_acc)
print("Logistic Regression Precision:", lr_prec)
print("Logistic Regression Recall:", lr_rec)
print("Logistic Regression F1 Score:", lr_f1)
```

```
Logistic Regression Accuracy: 0.8888888888888888
Logistic Regression Precision: 0.8857142857142857
Logistic Regression Recall: 0.9393939393939394
Logistic Regression F1 Score: 0.9117647058823529
```

```
[24]: import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix
import numpy as np
import itertools
```

```
[25]: # Define class labels
classes = ['Class A', 'Class B']
```

```
[26]: # Compute confusion matrix
cm = confusion_matrix(y_test, lr_preds)
```

```
[27]: # Plot confusion matrix
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion matrix')
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)
plt.xlabel('Predicted value')
plt.ylabel('True value')


# Add text to each cell
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], 'd'),
```

```

        horizontalalignment="center",
        color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.show()

```



0.0.2 Decision Tree

```

[28]: # Train and evaluate a Decision Tree model
dt_model = DecisionTreeClassifier(random_state=42)
dt_model.fit(X_train, y_train)

```

```

[28]: DecisionTreeClassifier(random_state=42)

```

```

[29]: dt_preds = dt_model.predict(X_test)
dt_acc = accuracy_score(y_test, dt_preds)
dt_prec = precision_score(y_test, dt_preds)
dt_rec = recall_score(y_test, dt_preds)
dt_f1 = f1_score(y_test, dt_preds)

```

```

[30]: print("Results of Decision Tree:")
print("Decision Tree Accuracy:", dt_acc)
print("Decision Tree Precision:", dt_prec)

```

```
print("Decision Tree Recall:", dt_rec)
print("Decision Tree F1 Score:", dt_f1)
```

Results of Decision Tree:

Decision Tree Accuracy: 0.7407407407407407

Decision Tree Precision: 0.8518518518518519


Decision Tree Recall: 0.696969696969697

Decision Tree F1 Score: 0.7666666666666667

```
[31]: from sklearn.metrics import plot_precision_recall_curve
import matplotlib.pyplot as plt
```

```
[32]: plot_precision_recall_curve(dt_model, X_test, y_test)
```

```
[32]: <sklearn.metrics._plot.precision_recall_curve.PrecisionRecallDisplay at
0x1ee4699fbb0>
```



0.0.3 Gradient Boosting

```
[33]: gb_model = GradientBoostingClassifier(random_state=42)
gb_model.fit(X_train, y_train)
```

```
[33]: GradientBoostingClassifier(random_state=42)
```

```
[34]: gb_preds = gb_model.predict(X_test)
```

```
[35]: gb_acc = accuracy_score(y_test, gb_preds)
      gb_prec = precision_score(y_test, gb_preds)
      gb_rec = recall_score(y_test, gb_preds)
      gb_f1 = f1_score(y_test, gb_preds)
```

```
[36]: print("Results of Gradient Boosting")
      print("Gradient Boosting Accuracy:", gb_acc)
      print("Gradient Boosting Precision:", gb_prec)
      print("Gradient Boosting Recall:", gb_rec)
      print("Gradient Boosting F1 Score:", gb_f1)
```

Results of Gradient Boosting

Gradient Boosting Accuracy: 0.7592592592592593

Gradient Boosting Precision: 0.7777777777777778

Gradient Boosting Recall: 0.8484848484848485

Gradient Boosting F1 Score: 0.8115942028985507