

Ionosphere dataset classification

March 24, 2023

```
[2]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
```

```
[3]: from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, f1_score, precision_score, \
    recall_score
```

```
[4]: import warnings
warnings.filterwarnings("ignore")
```

```
[5]: data = pd.read_csv('ionosphere.csv')
```

```
[6]: data.head(5)
```

```
[6]: @inputs Pulse1 Pulse3 Pulse4 Pulse5 Pulse6 Pulse7 Pulse8 Pulse9 \
0      1      0.995 -0.059  0.852  0.023  0.834 -0.377  1.000
1      1      1.000 -0.188  0.930 -0.362 -0.109 -0.936  1.000
2      1      1.000 -0.452  1.000  1.000  0.712 -1.000  0.000
3      1      1.000 -0.024  0.941  0.065  0.921 -0.233  0.772
4      1      0.023 -0.006 -0.099 -0.119 -0.008 -0.118  0.147

      Pulse10 Pulse11 ... Pulse26 Pulse27 Pulse28 Pulse29 Pulse30 \
0      0.038  0.852 ...  -0.512  0.411  -0.462  0.213  -0.341
1     -0.045  0.509 ...  -0.266 -0.205  -0.184 -0.190  -0.116
2      0.000  0.000 ...   0.907  0.516  1.000  1.000  -0.201
3     -0.164  0.528 ...  -0.652  0.133  -0.532  0.024  -0.622
4      0.066  0.038 ...  -0.015 -0.032  0.092 -0.079  0.007

      Pulse31 Pulse32 Pulse33 Pulse34 @outputs Class
0      0.423  -0.545   0.186  -0.453          g
1     -0.166  -0.063  -0.137  -0.024          b
2      0.257   1.000  -0.324   1.000          b
3     -0.057  -0.596  -0.046  -0.657          g
4      0.000   0.000   0.000   0.120          b
```

[5 rows x 34 columns]

```
[7]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 351 entries, 0 to 350
Data columns (total 34 columns):
#   Column                Non-Null Count  Dtype
---  -
0   @inputs Pulse1        351 non-null    int64
1   Pulse3                351 non-null    float64
2   Pulse4                351 non-null    float64
3   Pulse5                351 non-null    float64
4   Pulse6                351 non-null    float64
5   Pulse7                351 non-null    float64
6   Pulse8                351 non-null    float64
7   Pulse9                351 non-null    float64
8   Pulse10               351 non-null    float64
9   Pulse11               351 non-null    float64
10  Pulse12               351 non-null    float64
11  Pulse13               351 non-null    float64
12  Pulse14               351 non-null    float64
13  Pulse15               351 non-null    float64
14  Pulse16               351 non-null    float64
15  Pulse17               351 non-null    float64
16  Pulse18               351 non-null    float64
17  Pulse19               351 non-null    float64
18  Pulse20               351 non-null    float64
19  Pulse21               351 non-null    float64
20  Pulse22               351 non-null    float64
21  Pulse23               351 non-null    float64
22  Pulse24               351 non-null    float64
23  Pulse25               351 non-null    float64
24  Pulse26               351 non-null    float64
25  Pulse27               351 non-null    float64
26  Pulse28               351 non-null    float64
27  Pulse29               351 non-null    float64
28  Pulse30               351 non-null    float64
29  Pulse31               351 non-null    float64
30  Pulse32               351 non-null    float64
31  Pulse33               351 non-null    float64
32  Pulse34               351 non-null    float64
33  @outputs Class        351 non-null    object
dtypes: float64(32), int64(1), object(1)
memory usage: 93.4+ KB
```

```
[8]: data = data.rename(columns={"@inputs Pulse1": "Pulse1", "@outputs Class":
    ↪ "Class"})
```

```
[9]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 351 entries, 0 to 350
Data columns (total 34 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Pulse1      351 non-null   int64
1   Pulse3      351 non-null   float64
2   Pulse4      351 non-null   float64
3   Pulse5      351 non-null   float64
4   Pulse6      351 non-null   float64
5   Pulse7      351 non-null   float64
6   Pulse8      351 non-null   float64
7   Pulse9      351 non-null   float64
8   Pulse10     351 non-null   float64
9   Pulse11     351 non-null   float64
10  Pulse12     351 non-null   float64
11  Pulse13     351 non-null   float64
12  Pulse14     351 non-null   float64
13  Pulse15     351 non-null   float64
14  Pulse16     351 non-null   float64
15  Pulse17     351 non-null   float64
16  Pulse18     351 non-null   float64
17  Pulse19     351 non-null   float64
18  Pulse20     351 non-null   float64
19  Pulse21     351 non-null   float64
20  Pulse22     351 non-null   float64
21  Pulse23     351 non-null   float64
22  Pulse24     351 non-null   float64
23  Pulse25     351 non-null   float64
24  Pulse26     351 non-null   float64
25  Pulse27     351 non-null   float64
26  Pulse28     351 non-null   float64
27  Pulse29     351 non-null   float64
28  Pulse30     351 non-null   float64
29  Pulse31     351 non-null   float64
30  Pulse32     351 non-null   float64
31  Pulse33     351 non-null   float64
32  Pulse34     351 non-null   float64
33  Class       351 non-null   object
dtypes: float64(32), int64(1), object(1)
memory usage: 93.4+ KB
```

```
[10]: # Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(data.iloc[:, :-1], data.
↪iloc[:, -1], test_size=0.2, random_state=42)
```

0.0.1 Logistic Regression

```
[11]: # Train and evaluate a Logistic Regression model
lr_model = LogisticRegression(random_state=42)
```

```
[12]: lr_model.fit(X_train, y_train)
lr_preds = lr_model.predict(X_test)
```

```
[13]: lr_acc = accuracy_score(y_test, lr_preds)
lr_prec = precision_score(y_test, lr_preds, pos_label=' b')
lr_rec = recall_score(y_test, lr_preds, pos_label=' b')
lr_f1 = f1_score(y_test, lr_preds, pos_label=' b')
```

```
[14]: print("Logistic Regression Accuracy:", lr_acc)
print("Logistic Regression Precision:", lr_prec)
print("Logistic Regression Recall:", lr_rec)
print("Logistic Regression F1 Score:", lr_f1)
```

```
Logistic Regression Accuracy: 0.8028169014084507
Logistic Regression Precision: 0.8421052631578947
Logistic Regression Recall: 0.5925925925925926
Logistic Regression F1 Score: 0.6956521739130435
```

```
[15]: import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix
import numpy as np
import itertools
```

```
[16]: # Define class labels
classes = ['Class A', 'Class B']
```

```
[17]: # Compute confusion matrix
cm = confusion_matrix(y_test, lr_preds)
```

```
[18]: # Plot confusion matrix
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion matrix')
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)
plt.xlabel('Predicted value')
plt.ylabel('True value')


# Add text to each cell
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], 'd'),
```

```

        horizontalalignment="center",
        color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.show()

```



0.0.2 Decision Tree

```

[19]: # Train and evaluate a Decision Tree model
      dt_model = DecisionTreeClassifier()
      dt_model.fit(X_train, y_train)

```

```

[19]: DecisionTreeClassifier()

```

```

[20]: dt_preds = dt_model.predict(X_test)
      dt_acc = accuracy_score(y_test, dt_preds)
      dt_prec = precision_score(y_test, dt_preds, pos_label=' g')
      dt_rec = recall_score(y_test, dt_preds, pos_label=' g')
      dt_f1 = f1_score(y_test, dt_preds, pos_label=' g')

```

```

[21]: print("Results of Decision Tree:")
      print("Decision Tree Accuracy:", dt_acc)
      print("Decision Tree Precision:", dt_prec)

```

```
print("Decision Tree Recall:", dt_rec)
print("Decision Tree F1 Score:", dt_f1)
```

Results of Decision Tree:

Decision Tree Accuracy: 0.8591549295774648

Decision Tree Precision: 0.8863636363636364

Decision Tree Recall: 0.8863636363636364

Decision Tree F1 Score: 0.8863636363636365

```
[22]: import matplotlib.pyplot as plt
      from sklearn.metrics import roc_curve, auc
      from sklearn.preprocessing import LabelEncoder
```


```
[23]: # Calculate predicted probabilities for positive class
      dt_probs = dt_model.predict_proba(X_test)[:, 1]
```

```
[24]: y_true_binary = (y_test == ' g').astype(int)
```

```
[25]: # calculate FPR, TPR, and thresholds
      fpr, tpr, thresholds = roc_curve(y_true_binary, dt_probs, pos_label=1)
```

```
[26]: # Calculate AUC
      auc_dt = auc(fpr, tpr)
```

```
[27]: # Plot ROC curve
      plt.plot(fpr, tpr, label='Decision Tree (AUC = {:.2f})'.format(auc_dt))
      plt.plot([0, 1], [0, 1], linestyle='--', color='gray', label='Random Guess')
      plt.xlabel('False Positive Rate')
      plt.ylabel('True Positive Rate')
      plt.title('ROC Curve')
      plt.legend()
      plt.show()
```



0.0.3 Gradient Boosting

```
[28]: gb_model = GradientBoostingClassifier(random_state=42)
      gb_model.fit(X_train, y_train)
```

```
[28]: GradientBoostingClassifier(random_state=42)
```

```
[29]: gb_preds = gb_model.predict(X_test)
```

```
[30]: gb_acc = accuracy_score(y_test, gb_preds)
      gb_prec = precision_score(y_test, gb_preds, pos_label=' g')
      gb_rec = recall_score(y_test, gb_preds, pos_label=' g')
      gb_f1 = f1_score(y_test, gb_preds, pos_label=' g')
```

```
[36]: print("Results of Gradient Boosting")
      print("Gradient Boosting Accuracy:", gb_acc)
      print("Gradient Boosting Precision:", gb_prec)
      print("Gradient Boosting Recall:", gb_rec)
      print("Gradient Boosting F1 Score:", gb_f1)
```

```
Results of Gradient Boosting
Gradient Boosting Accuracy: 0.9014084507042254
Gradient Boosting Precision: 0.8936170212765957
```

Gradient Boosting Recall: 0.9545454545454546
Gradient Boosting F1 Score: 0.9230769230769231