

# Baseball

March 9, 2023

**0.0.1 In this Analysis Linear Regression, Neural Network, XG Boost and Logistic Regression method is used**

```
[1]: import pandas as pd
import numpy as np
from tensorflow.keras import layers
from tensorflow.keras import models
from sklearn.metrics import r2_score, accuracy_score

[2]: from sklearn.model_selection import KFold
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error, \
    confusion_matrix
import math
from sklearn.model_selection import train_test_split, GridSearchCV

[3]: # to ignore warnings
import warnings
warnings.filterwarnings("ignore")
```

**0.0.2 With the help of Pandas, read the ".csv" file and performing some task**

```
[4]: data1 = pd.read_csv("baseball.csv")

[5]: data1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 337 entries, 0 to 336
Data columns (total 17 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   @inputs Batting_average               337 non-null    float64
1   On-base_percentage                   337 non-null    float64
2   Runs                                 337 non-null    int64
3   Hits                                 337 non-null    int64
4   Doubles                              337 non-null    int64
5   Triples                              337 non-null    int64
6   HomeRuns                             337 non-null    int64
7   Runs_batted_in                       337 non-null    int64
```

```

8 Walks 337 non-null int64
9 Strike-Outs 337 non-null int64
10 Stolen_bases 337 non-null int64
11 Errors 337 non-null int64
12 Free_agency_eligibility 337 non-null int64
13 Free_agent 337 non-null int64
14 Arbitration_eligibility 337 non-null int64
15 Arbitration 337 non-null int64
16 @outputs Salary 337 non-null int64
dtypes: float64(2), int64(15)
memory usage: 44.9 KB

```

### 0.0.3 Renaming the columns for better understanding

```
[6]: data1.rename(columns = {'@inputs Batting_average': 'Batting_average', '@outputs_1 Salary': 'Salary'}, inplace = True)
```

```
[7]: data11 = data1
```

```
[8]: data1.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 337 entries, 0 to 336
Data columns (total 17 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Batting_average                       337 non-null    float64
1   On-base_percentage                   337 non-null    float64
2   Runs                                 337 non-null    int64
3   Hits                                 337 non-null    int64
4   Doubles                              337 non-null    int64
5   Triples                              337 non-null    int64
6   HomeRuns                             337 non-null    int64
7   Runs_batted_in                       337 non-null    int64
8   Walks                                337 non-null    int64
9   Strike-Outs                          337 non-null    int64
10  Stolen_bases                          337 non-null    int64
11  Errors                                337 non-null    int64
12  Free_agency_eligibility               337 non-null    int64
13  Free_agent                           337 non-null    int64
14  Arbitration_eligibility               337 non-null    int64
15  Arbitration                           337 non-null    int64
16  Salary                                337 non-null    int64
dtypes: float64(2), int64(15)
memory usage: 44.9 KB

```

```
[9]: data1.head(4)
```

```
[9]:
```

	Batting_average	On-base_percentage	Runs	Hits	Doubles	Triples	\
0	0.271	0.328	74	161	22	6	
1	0.264	0.318	24	48	7	0	
2	0.251	0.338	101	141	35	3	
3	0.224	0.274	28	94	21	1	

	HomeRuns	Runs_batted_in	Walks	Strike-Outs	Stolen_bases	Errors	\
0	12	58	49	133	23	17	
1	1	22	15	18	0	7	
2	32	105	71	104	34	6	
3	1	44	27	54	2	7	

	Free_agency_eligibility	Free_agent	Arbitration_eligibility	Arbitration	\
0	1	1	0	0	
1	0	0	0	0	
2	0	0	1	0	
3	1	1	0	0	

	Salary
0	109
1	160
2	2700
3	550

```
data1.tail(4)
```

```
[10]: data1.isnull().any()
```

```
[10]:
```

Batting_average	False
On-base_percentage	False
Runs	False
Hits	False
Doubles	False
Triples	False
HomeRuns	False
Runs_batted_in	False
Walks	False
Strike-Outs	False
Stolen_bases	False
Errors	False
Free_agency_eligibility	False
Free_agent	False
Arbitration_eligibility	False
Arbitration	False
Salary	False

```
dtype: bool
```

```
[11]: data1.isnull().sum()
```

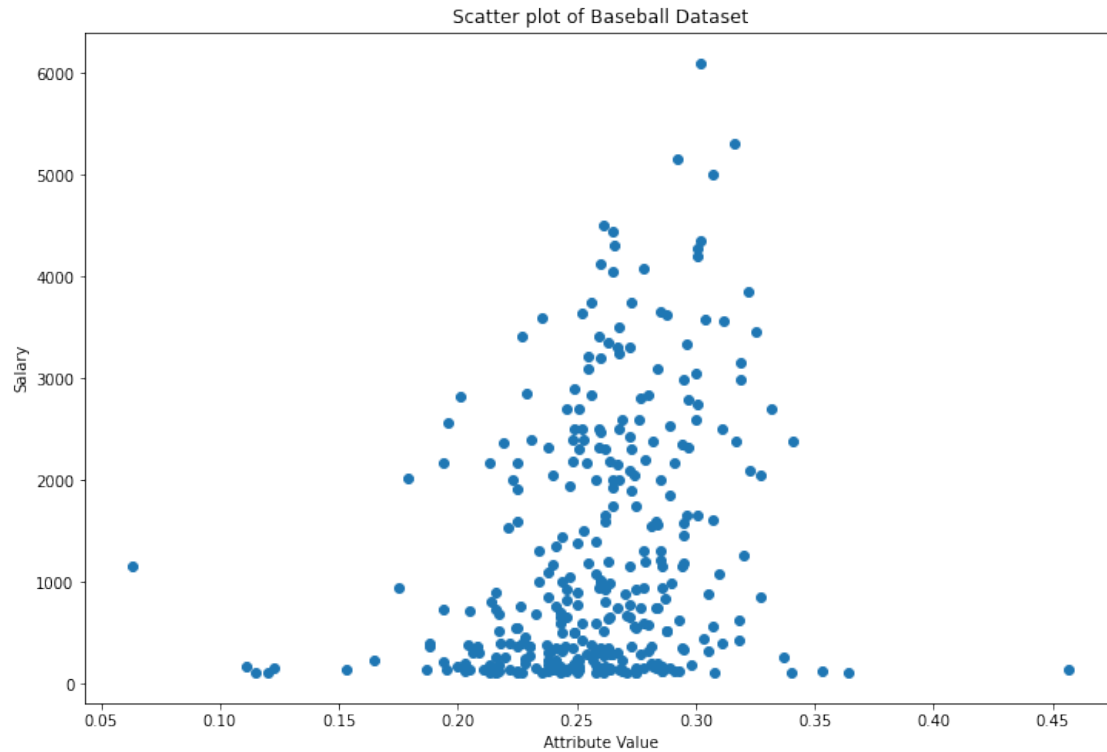
```
[11]: Batting_average      0
      On-base_percentage  0
      Runs                 0
      Hits                 0
      Doubles              0
      Triples              0
      HomeRuns             0
      Runs_batted_in       0
      Walks                0
      Strike-Outs          0
      Stolen_bases         0
      Errors               0
      Free_agency_eligibility 0
      Free_agent            0
      Arbitration_eligibility 0
      Arbitration           0
      Salary                0
      dtype: int64
```

```
[12]: import seaborn as sns
      import matplotlib.pyplot as plt
```

```
[13]: # Create the scatter plot
fig, ax = plt.subplots(figsize=(12,8))
ax.scatter(data1["Batting_average"], data1["Salary"])

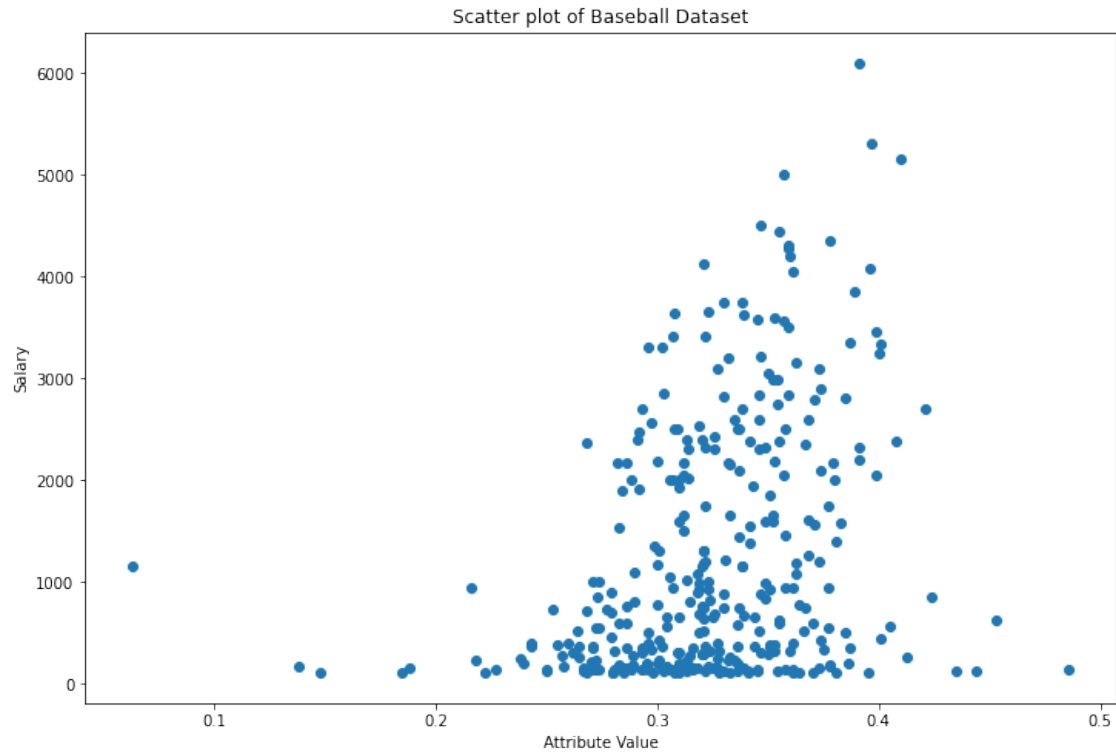
ax.set_xlabel("Attribute Value")
ax.set_ylabel("Salary")
ax.set_title("Scatter plot of Baseball Dataset")
plt.show()

#scatterplot is best used to detect Trends, Outlier detection
```




```
[14]: fig, ax = plt.subplots(figsize=(12,8))
      ax.scatter(data1["On-base_percentage"], data1["Salary"])

      ax.set_xlabel("Attribute Value")
      ax.set_ylabel("Salary")
      ax.set_title("Scatter plot of Baseball Dataset")
      plt.show()
```




```
[15]: fig, ax = plt.subplots(figsize=(12,8))
      ax.scatter(data1["Runs"], data1["Salary"])

      ax.set_xlabel("Attribute Value")
      ax.set_ylabel("Salary")
      ax.set_title("Scatter plot of Baseball Dataset")
      plt.show()
```



```
[16]: fig, ax = plt.subplots(figsize=(12,8))
      ax.scatter(data1["Hits"], data1["Salary"])


      ax.set_xlabel("Attribute Value")
      ax.set_ylabel("Salary")
      ax.set_title("Scatter plot of Baseball Dataset")
      plt.show()
```



```
[17]: fig, ax = plt.subplots(figsize=(12,8))
      ax.scatter(data1["Doubles"], data1["Salary"])
      ax.scatter(data1["Triples"], data1["Salary"])


      ax.set_xlabel("Attribute Value")
      ax.set_ylabel("Salary")
      ax.set_title("Scatter plot of Baseball Dataset")
      plt.show()
```






```
[18]: fig, ax = plt.subplots(figsize=(12,8))

ax.scatter(data1["HomeRuns"], data1["Salary"])
ax.scatter(data1["Runs_batted_in"], data1["Salary"])
ax.set_xlabel("Attribute Value")
ax.set_ylabel("Salary")
ax.set_title("Scatter plot of Baseball Dataset")
plt.show()
```




```
[19]: fig, ax = plt.subplots(figsize=(12,8))

ax.scatter(data1["Walks"], data1["Salary"])
ax.scatter(data1["Strike-Outs"], data1["Salary"])
ax.set_xlabel("Attribute Value")
ax.set_ylabel("Salary")
ax.set_title("Scatter plot of Baseball Dataset")
plt.show()
```



```
[20]: fig, ax = plt.subplots(figsize=(12,8))

ax.scatter(data1["Stolen_bases"], data1["Salary"])
ax.scatter(data1["Errors"], data1["Salary"])
ax.set_xlabel("Attribute Value")
ax.set_ylabel("Salary")
ax.set_title("Scatter plot of Baseball Dataset")
plt.show()
```



```
[21]: # Split the dataset into training and testing sets
X = data1.drop("Salary", axis=1) # Independent variables
y = data1["Salary"] # Dependent variable
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

#### 0.0.4 Building the neural network model

```
[22]: def build_model():
    model = models.Sequential()
    model.add(layers.Dense(64, activation='relu', input_shape=(X.shape[1],)))
    model.add(layers.Dense(64, activation='relu'))
    model.add(layers.Dense(1))
    model.compile(optimizer='rmsprop', loss='mse', metrics=['mae'])
    return model
```

#### 0.0.5 Define and perform K-fold cross validation

```
[23]: kf = KFold(n_splits=5)
```

```
[24]: # Perform k-fold cross validation
for train_index, test_index in kf.split(X):
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
```

```
y_train, y_test = y.iloc[train_index], y.iloc[test_index]
```

### 0.0.6 With neural network

```
[25]: # Standardize the training and testing data
X_train_mean = X_train.mean()
X_train_std = X_train.std()
X_train = (X_train - X_train_mean) / X_train_std
X_test = (X_test - X_train_mean) / X_train_std
```

```
[26]: # Build the model
model = build_model()
```

```
[27]: history = model.fit(X_train, y_train, epochs=30 , batch_size=5, verbose=0)
```

```
[28]: r2_scores = []
```

```
[29]: # Predict the target values for the testing data
y_pred = model.predict(X_test).flatten()
```

```
3/3 [=====] - 0s 977us/step
```

```
[30]: r2 = r2_score(y_test, y_pred)
r2_scores.append(r2)
```

```
[31]: print('Coefficient of determination R-squared (R2):', np.mean(r2_scores))
```

```
Coefficient of determination R-squared (R2): 0.5381060070607421
```

### 0.0.7 Initialize the Logistic regression and fit the model

```
[32]: # Create a binary target variable
data1["HighSalary"] = (data1["Salary"] > data1["Salary"].median()).astype(int)
```

```
[33]: # Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(data1[["Batting_average",
↪ "On-base_percentage", "Runs", "Hits", "Doubles", "Triples", "HomeRuns",
↪ "Runs_batted_in", "Walks", "Strike-Outs", "Stolen_bases", "Errors"]],
↪ data1["HighSalary"], test_size=0.3, random_state=42)
```

```
[34]: model = LogisticRegression()
```

```
[35]: model.fit(X_train, y_train)
```

```
[35]: LogisticRegression()
```


```
[36]: y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
```

```
[37]: print("Accuracy: {:.2f}%".format(accuracy*100))
```

Accuracy: 75.49%

```
[38]: # Create a confusion matrix
cm = confusion_matrix(y_test, y_pred, labels=[0, 1])

# Plot the confusion matrix
sns.heatmap(cm, annot=True, cmap="Blues", fmt="g", xticklabels=["Low", "High"],
            yticklabels=["Low", "High"])
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title("Confusion Matrix")
plt.show()
```



#### 0.0.8 Initialize the linear regression model and fit the model

```
[39]: X_train, X_test, y_train, y_test = train_test_split(data11.drop("Salary",
            axis=1), data11["Salary"], test_size=0.2, random_state=42)
```

```
[40]: model = LinearRegression()
```

```
[41]: model.fit(X_train, y_train)
```

```
[41]: LinearRegression()
```

```
[42]: # Make predictions on the test data
y_pred = model.predict(X_test)
```

```
[43]: # Calculating the R2 Score
r2 = r2_score(y_test, y_pred)
print("Coefficient of determination R-squared (R2): ", r2)
```

Coefficient of determination R-squared (R2): 0.7858719726763712

### 0.0.9 By using XG Boost algorithm

```
[44]: import xgboost as xgb
```

```
[45]: # Define the XGBoost model
# xgb_model = xgb.XGBRegressor(objective='reg:squarederror')
xgb_model = xgb.XGBRegressor()
```

```
[46]: param_grid = {
    'learning_rate': [0.01, 0.05, 0.1],
    'max_depth': [3, 5, 7],
    'n_estimators': [50, 100, 200]
}
```

```
[47]: # Setting up cross-validation
xgb_cv = GridSearchCV(xgb_model, param_grid, cv=5)
```

```
[48]: xgb_model.fit(X_train, y_train)
```

```
[48]: XGBRegressor(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=None, early_stopping_rounds=None,
    enable_categorical=False, eval_metric=None, feature_types=None,
    gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
    interaction_constraints=None, learning_rate=None, max_bin=None,
    max_cat_threshold=None, max_cat_to_onehot=None,
    max_delta_step=None, max_depth=None, max_leaves=None,
    min_child_weight=None, missing=nan, monotone_constraints=None,
    n_estimators=100, n_jobs=None, num_parallel_tree=None,
    predictor=None, random_state=None, ...)
```

```
[49]: # Predict the target values for the testing data
y_pred = xgb_model.predict(X_test)
```

```
[50]: r2_scores = []
```

```
[51]: # Evaluate the model on the testing data
r2 = r2_score(y_test, y_pred)
```

```
[52]: r2_scores.append(r2)
```

```
[53]: # Compute the mean evaluation metrics over all the folds
print('Coefficient of determination R-squared (R2):', np.mean(r2_scores))
```


Coefficient of determination R-squared (R2): 0.7914602652460396

```
[54]: # Selecting the relevant columns for the regression
x = data1['Batting_average']
y = data1['Salary']

#enlarging the figure for better visualization
fig, ax = plt.subplots(figsize=(15, 10))

# Fit a regression model
sns.regplot(x, y, ax=ax)

# Show the plot
plt.show()
```



#### By Sushan Shankar