

# DEE

March 18, 2023

```
[1]: import pandas as pd
import numpy as np
from tensorflow.keras import layers
from tensorflow.keras import models
from sklearn.metrics import r2_score

[2]: from sklearn.model_selection import cross_val_score, KFold
from sklearn.linear_model import LinearRegression
from sklearn.neural_network import MLPRegressor
from xgboost import XGBRegressor

[3]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from keras.models import Sequential
from keras.layers import Dense, Dropout

[4]: # to ignore warnings
import warnings
warnings.filterwarnings("ignore")
```

## 0.0.1 With the help of Pandas, read the “.csv” file and performing some task

```
[5]: data1 = pd.read_csv("dee.csv")
```

```
[6]: data1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 365 entries, 0 to 364
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Hydroelectric    365 non-null   float64
1   Nuclear          365 non-null   int64
2   Coal             365 non-null   float64
3   Fuel             365 non-null   float64
4   Gas              365 non-null   float64
5   Special          365 non-null   int64
```

```
6    Consume          365 non-null    float64
dtypes: float64(5), int64(2)
memory usage: 20.1 KB
```

```
[7]: data1.head(4)
```

```
[7]:   Hydroelectric  Nuclear    Coal    Fuel    Gas  Special  Consume
0      179183.0    175973  78429.1  4680.73  8117.13     8023  1.732800
1      206035.0    186774  79129.5  4342.43  5715.18     8159  1.583500
2      198435.0    180633  64465.2  4566.84    0.00     8215  1.505310
3      187029.0    171382  51913.4  5342.54    0.00     8346  0.955205
```

```
[8]: data1.tail(4)
```

```
[8]:   Hydroelectric  Nuclear    Coal    Fuel    Gas  Special  Consume
361      139186.0    179138  122708.0   8628.45  6884.80    12382  2.01027
362      162721.0    174891  182417.0  16296.30  31574.60    15540  2.12432
363      109797.0    172697   96596.3   5803.86   3245.38    11801  2.06576
364      130802.0    179072  127273.0   5733.01   5344.81    12929  1.43698
```

```
[9]: data1.isnull().any()
```

```
[9]: Hydroelectric    False
Nuclear            False
Coal               False
Fuel              False
Gas               False
Special           False
Consume           False
dtype: bool
```


```
[10]: data1.isnull().sum()
```

```
[10]: Hydroelectric    0
Nuclear             0
Coal                0
Fuel               0
Gas                0
Special            0
Consume            0
dtype: int64
```

```
[11]: import seaborn as sns
import matplotlib.pyplot as plt
```

```
[12]: # Create a boxplot of the age variable
plt.boxplot(data1["Hydroelectric"])
plt.xlabel("Hydroelectric")
```

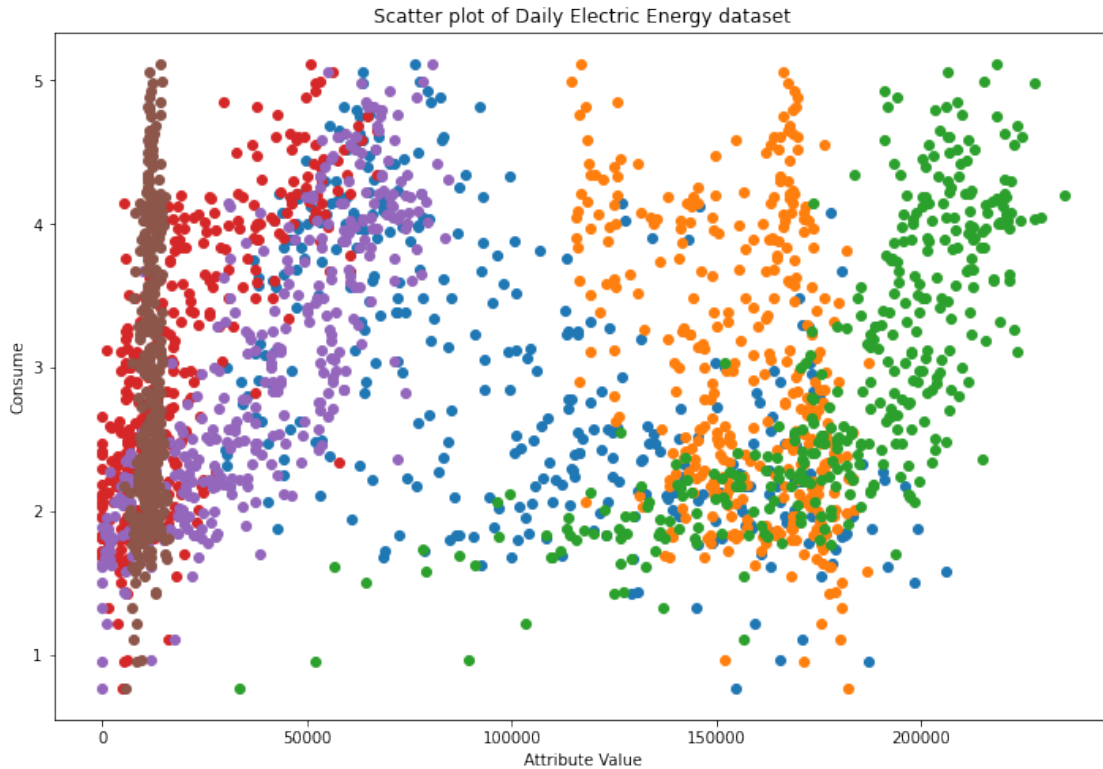
```
plt.show()
```



```
[13]: fig, ax = plt.subplots(figsize=(12,8))

ax.scatter(data1["Hydroelectric"], data1["Consume"])
ax.scatter(data1["Nuclear"], data1["Consume"])
ax.scatter(data1["Coal"], data1["Consume"])
ax.scatter(data1["Fuel"], data1["Consume"])
ax.scatter(data1["Gas"], data1["Consume"])
ax.scatter(data1["Special"], data1["Consume"])

ax.set_xlabel("Attribute Value")
ax.set_ylabel("Consume")
ax.set_title("Scatter plot of Daily Electric Energy dataset")
plt.show()
```



```
[14]: X = data1.drop('Consume', axis=1).values
      y = data1['Consume'].values
```

### 0.0.2 Create a linear regression model

```
[15]: linear_model = LinearRegression()
```

```
[16]: # Define the cross-validation method
      cv = KFold(n_splits=5, shuffle=True, random_state=42)
```

```
[17]: # Evaluate the model using 5-fold cross-validation
      linear_scores = cross_val_score(linear_model, X, y, cv=cv,
      ↪scoring='neg_mean_squared_error')
```

```
[18]: # Calculate the evaluation metrics from the scores
      rmse = np.sqrt(-linear_scores.mean())
      mse = -linear_scores.mean()
```

```
[19]: mae = -cross_val_score(linear_model, X, y, cv=cv,
      ↪scoring='neg_mean_absolute_error').mean()
      r2 = cross_val_score(linear_model, X, y, cv=cv, scoring='r2').mean()
```

```
[20]: # Print the evaluation metrics for the Linear Regression model
print('Linear Regression Model:')
print('RMSE:', rmse)
print('MSE:', mse)
print('MAE:', mae)
print('R-squared:', r2)
```

```
Linear Regression Model:
RMSE: 0.41281996865485276
MSE: 0.17042032652019362
MAE: 0.31678813633724523
R-squared: 0.8159626900027581
```

```
[21]: linear_model.fit(X, y)
```

```
[21]: LinearRegression()
```

```
[22]: new_data = [[27881.8, 114760.0, 33537.0, 0.0, 0.0, 5307.0]]
prediction = linear_model.predict(new_data)

# Print the predicted value
print('Predicted value of Consume:', prediction[0])
```

```
Predicted value of Consume: 1.6479062344769242
```

### 0.0.3 Create an Artificial Neural Network model

```
[23]: # Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

```
[24]: # Scale the features using standard scaler
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
[25]: # Create an Artificial Neural Network model
model = Sequential()
model.add(Dense(units=16, activation='relu', input_dim=X_train.shape[1]))
model.add(Dropout(rate=0.2))
model.add(Dense(units=8, activation='relu'))
model.add(Dropout(rate=0.2))
model.add(Dense(units=1))
model.compile(optimizer='adam', loss='mean_squared_error')
```

```
[26]: # Train the model
model.fit(X_train, y_train, epochs=100, batch_size=64, validation_split=0.2)
```

Epoch 1/100  
4/4 [=====] - 1s 83ms/step - loss: 11.9464 - val\_loss: 9.0904

Epoch 2/100  
4/4 [=====] - 0s 11ms/step - loss: 11.6631 - val\_loss: 8.8493

Epoch 3/100  
4/4 [=====] - 0s 11ms/step - loss: 11.5139 - val\_loss: 8.6185

Epoch 4/100  
4/4 [=====] - 0s 12ms/step - loss: 11.1444 - val\_loss: 8.4050

Epoch 5/100  
4/4 [=====] - 0s 11ms/step - loss: 11.0753 - val\_loss: 8.1990

Epoch 6/100  
4/4 [=====] - 0s 12ms/step - loss: 10.5356 - val\_loss: 8.0008

Epoch 7/100  
4/4 [=====] - 0s 11ms/step - loss: 10.2600 - val\_loss: 7.8061

Epoch 8/100  
4/4 [=====] - 0s 12ms/step - loss: 10.2784 - val\_loss: 7.6097

Epoch 9/100  
4/4 [=====] - 0s 12ms/step - loss: 9.9282 - val\_loss: 7.4146

Epoch 10/100  
4/4 [=====] - 0s 11ms/step - loss: 9.5919 - val\_loss: 7.2209

Epoch 11/100  
4/4 [=====] - 0s 11ms/step - loss: 9.4223 - val\_loss: 7.0205

Epoch 12/100  
4/4 [=====] - 0s 10ms/step - loss: 9.2295 - val\_loss: 6.8169

Epoch 13/100  
4/4 [=====] - 0s 11ms/step - loss: 9.0275 - val\_loss: 6.6047

Epoch 14/100  
4/4 [=====] - 0s 12ms/step - loss: 8.7099 - val\_loss: 6.3796

Epoch 15/100  
4/4 [=====] - 0s 12ms/step - loss: 8.3753 - val\_loss: 6.1463

Epoch 16/100  
4/4 [=====] - 0s 11ms/step - loss: 8.2154 - val\_loss: 5.9046

Epoch 17/100  
4/4 [=====] - 0s 11ms/step - loss: 7.9283 - val\_loss: 5.6511  
Epoch 18/100  
4/4 [=====] - 0s 13ms/step - loss: 7.4927 - val\_loss: 5.3902  
Epoch 19/100  
4/4 [=====] - 0s 12ms/step - loss: 7.0290 - val\_loss: 5.1229  
Epoch 20/100  
4/4 [=====] - 0s 12ms/step - loss: 6.9064 - val\_loss: 4.8531  
Epoch 21/100  
4/4 [=====] - 0s 12ms/step - loss: 6.5768 - val\_loss: 4.5813  
Epoch 22/100  
4/4 [=====] - 0s 15ms/step - loss: 5.9098 - val\_loss: 4.2976  
Epoch 23/100  
4/4 [=====] - 0s 11ms/step - loss: 5.8229 - val\_loss: 4.0132  
Epoch 24/100  
4/4 [=====] - 0s 11ms/step - loss: 5.5866 - val\_loss: 3.7301  
Epoch 25/100  
4/4 [=====] - 0s 10ms/step - loss: 5.2033 - val\_loss: 3.4552  
Epoch 26/100  
4/4 [=====] - 0s 10ms/step - loss: 4.7514 - val\_loss: 3.1886  
Epoch 27/100  
4/4 [=====] - 0s 12ms/step - loss: 4.3201 - val\_loss: 2.9281  
Epoch 28/100  
4/4 [=====] - 0s 11ms/step - loss: 4.3296 - val\_loss: 2.6822  
Epoch 29/100  
4/4 [=====] - 0s 12ms/step - loss: 3.9671 - val\_loss: 2.4462  
Epoch 30/100  
4/4 [=====] - 0s 11ms/step - loss: 3.5596 - val\_loss: 2.2225  
Epoch 31/100  
4/4 [=====] - 0s 11ms/step - loss: 3.2157 - val\_loss: 2.0067  
Epoch 32/100  
4/4 [=====] - 0s 10ms/step - loss: 3.0046 - val\_loss: 1.8048

Epoch 33/100  
4/4 [=====] - 0s 10ms/step - loss: 2.5191 - val\_loss: 1.6214  
Epoch 34/100  
4/4 [=====] - 0s 11ms/step - loss: 2.4460 - val\_loss: 1.4536  
Epoch 35/100  
4/4 [=====] - 0s 11ms/step - loss: 2.5960 - val\_loss: 1.3020  
Epoch 36/100  
4/4 [=====] - 0s 13ms/step - loss: 2.5010 - val\_loss: 1.1621  
Epoch 37/100  
4/4 [=====] - 0s 10ms/step - loss: 1.8669 - val\_loss: 1.0408  
Epoch 38/100  
4/4 [=====] - 0s 10ms/step - loss: 2.0692 - val\_loss: 0.9352  
Epoch 39/100  
4/4 [=====] - 0s 11ms/step - loss: 1.7135 - val\_loss: 0.8413  
Epoch 40/100  
4/4 [=====] - 0s 11ms/step - loss: 2.0060 - val\_loss: 0.7643  
Epoch 41/100  
4/4 [=====] - 0s 11ms/step - loss: 1.5310 - val\_loss: 0.7001  
Epoch 42/100  
4/4 [=====] - 0s 10ms/step - loss: 1.7147 - val\_loss: 0.6487  
Epoch 43/100  
4/4 [=====] - 0s 11ms/step - loss: 1.6232 - val\_loss: 0.6057  
Epoch 44/100  
4/4 [=====] - 0s 11ms/step - loss: 1.7053 - val\_loss: 0.5674  
Epoch 45/100  
4/4 [=====] - 0s 10ms/step - loss: 1.5380 - val\_loss: 0.5368  
Epoch 46/100  
4/4 [=====] - 0s 11ms/step - loss: 1.5218 - val\_loss: 0.5126  
Epoch 47/100  
4/4 [=====] - 0s 12ms/step - loss: 1.6862 - val\_loss: 0.4921  
Epoch 48/100  
4/4 [=====] - 0s 12ms/step - loss: 1.5758 - val\_loss: 0.4742



Epoch 49/100  
4/4 [=====] - 0s 11ms/step - loss: 1.5785 - val\_loss: 0.4597  
Epoch 50/100  
4/4 [=====] - 0s 11ms/step - loss: 1.4930 - val\_loss: 0.4472  
Epoch 51/100  
4/4 [=====] - 0s 10ms/step - loss: 1.4261 - val\_loss: 0.4370  
Epoch 52/100  
4/4 [=====] - 0s 12ms/step - loss: 1.4445 - val\_loss: 0.4272  
Epoch 53/100  
4/4 [=====] - 0s 11ms/step - loss: 1.4492 - val\_loss: 0.4213  
Epoch 54/100  
4/4 [=====] - 0s 11ms/step - loss: 1.3140 - val\_loss: 0.4167  
Epoch 55/100  
4/4 [=====] - 0s 10ms/step - loss: 1.4822 - val\_loss: 0.4090  
Epoch 56/100  
4/4 [=====] - 0s 11ms/step - loss: 1.4581 - val\_loss: 0.4028  
Epoch 57/100  
4/4 [=====] - 0s 11ms/step - loss: 1.3464 - val\_loss: 0.3944  
Epoch 58/100  
4/4 [=====] - 0s 10ms/step - loss: 1.4401 - val\_loss: 0.3869  
Epoch 59/100  
4/4 [=====] - 0s 10ms/step - loss: 1.3851 - val\_loss: 0.3796  
Epoch 60/100  
4/4 [=====] - 0s 10ms/step - loss: 1.3708 - val\_loss: 0.3700  
Epoch 61/100  
4/4 [=====] - 0s 10ms/step - loss: 1.6417 - val\_loss: 0.3617  
Epoch 62/100  
4/4 [=====] - 0s 10ms/step - loss: 1.3589 - val\_loss: 0.3556  
Epoch 63/100  
4/4 [=====] - 0s 11ms/step - loss: 1.2982 - val\_loss: 0.3515  
Epoch 64/100  
4/4 [=====] - 0s 10ms/step - loss: 1.3851 - val\_loss: 0.3482

Epoch 65/100  
4/4 [=====] - 0s 10ms/step - loss: 1.5463 - val\_loss:  
0.3428  
Epoch 66/100  
4/4 [=====] - 0s 11ms/step - loss: 1.4726 - val\_loss:  
0.3401  
Epoch 67/100  
4/4 [=====] - 0s 11ms/step - loss: 1.0728 - val\_loss:  
0.3370  
Epoch 68/100  
4/4 [=====] - 0s 12ms/step - loss: 1.2520 - val\_loss:  
0.3346  
Epoch 69/100  
4/4 [=====] - 0s 10ms/step - loss: 1.2275 - val\_loss:  
0.3327  
Epoch 70/100  
4/4 [=====] - 0s 11ms/step - loss: 1.4167 - val\_loss:  
0.3303  
Epoch 71/100  
4/4 [=====] - 0s 10ms/step - loss: 1.1346 - val\_loss:  
0.3295  
Epoch 72/100  
4/4 [=====] - 0s 11ms/step - loss: 1.3737 - val\_loss:  
0.3309  
Epoch 73/100  
4/4 [=====] - 0s 10ms/step - loss: 1.1714 - val\_loss:  
0.3358  
Epoch 74/100  
4/4 [=====] - 0s 11ms/step - loss: 1.2978 - val\_loss:  
0.3400  
Epoch 75/100  
4/4 [=====] - 0s 13ms/step - loss: 1.4047 - val\_loss:  
0.3437  
Epoch 76/100  
4/4 [=====] - 0s 13ms/step - loss: 1.0385 - val\_loss:  
0.3455  
Epoch 77/100  
4/4 [=====] - 0s 12ms/step - loss: 1.3871 - val\_loss:  
0.3479  
Epoch 78/100  
4/4 [=====] - 0s 13ms/step - loss: 1.2825 - val\_loss:  
0.3457  
Epoch 79/100  
4/4 [=====] - 0s 21ms/step - loss: 1.2230 - val\_loss:  
0.3405  
Epoch 80/100  
4/4 [=====] - 0s 12ms/step - loss: 1.1306 - val\_loss:  
0.3368

Epoch 81/100  
4/4 [=====] - 0s 11ms/step - loss: 0.9573 - val\_loss: 0.3330  
Epoch 82/100  
4/4 [=====] - 0s 11ms/step - loss: 1.0028 - val\_loss: 0.3293  
Epoch 83/100  
4/4 [=====] - 0s 13ms/step - loss: 1.2515 - val\_loss: 0.3263  
Epoch 84/100  
4/4 [=====] - 0s 11ms/step - loss: 1.1738 - val\_loss: 0.3234  
Epoch 85/100  
4/4 [=====] - 0s 11ms/step - loss: 1.2485 - val\_loss: 0.3217  
Epoch 86/100  
4/4 [=====] - 0s 11ms/step - loss: 1.4123 - val\_loss: 0.3206  
Epoch 87/100  
4/4 [=====] - 0s 10ms/step - loss: 1.0697 - val\_loss: 0.3158  
Epoch 88/100  
4/4 [=====] - 0s 10ms/step - loss: 1.1360 - val\_loss: 0.3105  
Epoch 89/100  
4/4 [=====] - 0s 10ms/step - loss: 1.0149 - val\_loss: 0.3055  
Epoch 90/100  
4/4 [=====] - 0s 10ms/step - loss: 1.0986 - val\_loss: 0.3032  
Epoch 91/100  
4/4 [=====] - 0s 11ms/step - loss: 1.0665 - val\_loss: 0.3019  
Epoch 92/100  
4/4 [=====] - 0s 12ms/step - loss: 1.1981 - val\_loss: 0.3013  
Epoch 93/100  
4/4 [=====] - 0s 11ms/step - loss: 0.8736 - val\_loss: 0.2972  
Epoch 94/100  
4/4 [=====] - 0s 11ms/step - loss: 1.0582 - val\_loss: 0.2932  
Epoch 95/100  
4/4 [=====] - 0s 13ms/step - loss: 1.1480 - val\_loss: 0.2907  
Epoch 96/100  
4/4 [=====] - 0s 10ms/step - loss: 1.2125 - val\_loss: 0.2896

```
Epoch 97/100
4/4 [=====] - 0s 10ms/step - loss: 1.0414 - val_loss: 0.2879
Epoch 98/100
4/4 [=====] - 0s 11ms/step - loss: 1.2164 - val_loss: 0.2870
Epoch 99/100
4/4 [=====] - 0s 11ms/step - loss: 1.1396 - val_loss: 0.2891
Epoch 100/100
4/4 [=====] - 0s 10ms/step - loss: 1.1171 - val_loss: 0.2903
```

[26]: <keras.callbacks.History at 0x2b6f7bb5f30>

```
[27]: # Make predictions on the test set
y_pred = model.predict(X_test)
```

```
3/3 [=====] - 0s 1ms/step
```

```
[28]: # Calculate the evaluation metrics
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
mse = mean_squared_error(y_test, y_pred)
```

```
[29]: mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
```

```
[30]: # Print the evaluation metrics for the Artificial Neural Network model
print('Artificial Neural Network Model:')
print('RMSE:', rmse)
print('MSE:', mse)
print('MAE:', mae)
print('R-squared:', r2)
```

```
Artificial Neural Network Model:
RMSE: 0.5472146718518526
MSE: 0.2994438970899307
MAE: 0.4252916363822205
R-squared: 0.653178329338943
```

#### 0.0.4 Create an XGBoost Regression model

```
[31]: xgb_model = XGBRegressor(objective='reg:squarederror', random_state=42)
```

```
[32]: # Evaluate the model using 5-fold cross-validation
xgb_scores = cross_val_score(xgb_model, X, y, cv=cv,
    ↪scoring='neg_mean_squared_error')
```

```
[33]: # Calculate the evaluation metrics from the scores
xgb_rmse = np.sqrt(-xgb_scores.mean())
xgb_mse = -xgb_scores.mean()

[34]: xgb_mae = -cross_val_score(xgb_model, X, y, cv=cv,
    ↪scoring='neg_mean_absolute_error').mean()
xgb_r2 = cross_val_score(xgb_model, X, y, cv=cv, scoring='r2').mean()

[35]: # Print the evaluation metrics for the XGBoost Regression model
print('XGBoost Regression Model:')
print('RMSE:', xgb_rmse)
print('MSE:', xgb_mse)
print('MAE:', xgb_mae)
print('R-squared:', xgb_r2)
```


```
XGBoost Regression Model:
RMSE: 0.4554426365691612
MSE: 0.20742799520506905
MAE: 0.3363558875635225
R-squared: 0.7753485901175845
```

### 0.0.5 Visualization

```
[36]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

[37]: # Compute the correlation matrix
corr = data1.corr()

# Plot the correlation matrix as a heatmap
sns.heatmap(corr, cmap='coolwarm', annot=True, fmt='.2f')
plt.title('Correlation Matrix for DEE Dataset')
plt.show()
```



```
[38]: # Plot the distribution of the target variable "Consume"
plt.hist(data1['Consume'], bins=20)
plt.xlabel('Daily Average Price of Electricity Energy (TkWhe)')
plt.ylabel('Frequency')
plt.title('Distribution of Daily Average Price of Electricity Energy')
plt.show()
```

