

# friedman

March 21, 2023

```
[1]: import pandas as pd
import numpy as np
from tensorflow.keras import layers
from tensorflow.keras import models
from sklearn.metrics import r2_score

[2]: from sklearn.model_selection import cross_val_score, KFold
from sklearn.linear_model import LinearRegression
from sklearn.neural_network import MLPRegressor
from xgboost import XGBRegressor

[3]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from keras.models import Sequential
from keras.layers import Dense, Dropout

[4]: # to ignore warnings
import warnings
warnings.filterwarnings("ignore")
```

## 0.0.1 With the help of Pandas, read the ".csv" file and performing some task

```
[5]: data1 = pd.read_csv("friedman.csv")

[6]: data1.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1202 entries, 0 to 1201
Data columns (total 6 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   @inputs Input1        1200 non-null   float64
 1   Input2                1200 non-null   float64
 2   Input3                1200 non-null   float64
 3   Input4                1200 non-null   float64
 4   Input5                1200 non-null   float64
 5   @outputs Output       1200 non-null   float64
```

```
dtypes: float64(6)
memory usage: 56.5 KB
```

```
[7]: data1.rename(columns = {'@inputs Input1': 'Input1', '@outputs Output': 'Output'},  
    ↪ inplace = True)
```

```
[8]: data1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1202 entries, 0 to 1201
Data columns (total 6 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0   Input1  1200 non-null    float64
 1   Input2  1200 non-null    float64
 2   Input3  1200 non-null    float64
 3   Input4  1200 non-null    float64
 4   Input5  1200 non-null    float64
 5   Output  1200 non-null    float64
dtypes: float64(6)
memory usage: 56.5 KB
```

```
[9]: data1.head(4)
```

```
[9]:
```

	Input1	Input2	Input3	Input4	Input5	Output
0	NaN	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	NaN	NaN
2	0.696482	0.358437	0.425834	0.330314	0.222491	11.094962
3	0.590390	0.430675	0.869042	0.070912	0.634303	13.229209

```
[10]: data1.tail(4)
```

```
[10]:
```

	Input1	Input2	Input3	Input4	Input5	Output
1198	0.348151	0.406174	0.243864	0.201591	0.040682	7.071847
1199	0.839787	0.759799	0.193053	0.187603	0.658195	14.336152
1200	0.017182	0.959536	0.815701	0.213163	0.681054	8.318943
1201	0.347079	0.870634	0.706439	0.147060	0.489136	13.031322

```
[11]: data1.isnull().any()
```

```
[11]: Input1    True
      Input2    True
      Input3    True
      Input4    True
      Input5    True
      Output    True
      dtype: bool
```

```
[12]: data1.isnull().sum()
```

```
[12]: Input1    2  
      Input2    2  
      Input3    2  
      Input4    2  
      Input5    2  
      Output    2  
      dtype: int64
```

```
[13]: null_cols = data1.isnull().sum()
```

```
[14]: # Loop through each input column and impute with mean  
      for col in data1.columns[:-1]:  
          mean_val = data1[col].mean()  
          data1[col].fillna(mean_val, inplace=True)
```

```
[15]: data1.isnull().sum()
```

```
[15]: Input1    0  
      Input2    0  
      Input3    0  
      Input4    0  
      Input5    0  
      Output    2  
      dtype: int64
```

```
[16]: # Impute null values in output column with mean  
      mean_val1 = data1['Output'].mean()  
      data1['Output'].fillna(mean_val1, inplace=True)
```

```
[17]: data1.isnull().sum()
```

```
[17]: Input1    0  
      Input2    0  
      Input3    0  
      Input4    0  
      Input5    0  
      Output    0  
      dtype: int64
```

```
[18]: import matplotlib.pyplot as plt
```

```
[19]: fig, ax = plt.subplots(figsize=(12,8))  
  
      ax.scatter(data1["Input1"], data1["Output"])  
      ax.scatter(data1["Input2"], data1["Output"])  
      ax.scatter(data1["Input3"], data1["Output"])
```


```

ax.scatter(data1["Input4"], data1["Output"])
ax.scatter(data1["Input5"], data1["Output"])

ax.set_xlabel("Attribute Value")

ax.set_title("Scatter plot of Elevator Dataset")
plt.show()

```



```

[20]: X = data1.drop('Output', axis=1).values
      y = data1['Output'].values

```

### 0.0.2 Create a linear regression model

```

[21]: linear_model = LinearRegression()

```

```

[22]: # Define the cross-validation method
      cv = KFold(n_splits=5, shuffle=True, random_state=42)

```

```

[23]: # Evaluate the model using 5-fold cross-validation
      linear_scores = cross_val_score(linear_model, X, y, cv=cv,
      ↪scoring='neg_mean_squared_error')

```

```
[24]: # Calculate the evaluation metrics from the scores
rmse = np.sqrt(-linear_scores.mean())
mse = -linear_scores.mean()
```

```
[25]: mae = -cross_val_score(linear_model, X, y, cv=cv,
    ↪scoring='neg_mean_absolute_error').mean()
r2 = cross_val_score(linear_model, X, y, cv=cv, scoring='r2').mean()
```

```
[26]: #the Linear Regression model
print('Linear Regression Model:')
print('RMSE:', rmse)
print('MSE:', mse)
print('MAE:', mae)
print('R-squared:', r2)
```

```
Linear Regression Model:
RMSE: 2.6962932503590205
MSE: 7.269997291931611
MAE: 2.0845673089752363
R-squared: 0.7281145447623009
```

### 0.0.3 Create an Artificial Neural Network model

```
[27]: # Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)
```

```
[28]: # Scale the features using standard scaler
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
[29]: # Create an Artificial Neural Network model
model = Sequential()
model.add(Dense(units=16, activation='relu', input_dim=X_train.shape[1]))
model.add(Dropout(rate=0.2))
model.add(Dense(units=8, activation='relu'))
model.add(Dropout(rate=0.2))
model.add(Dense(units=1))
model.compile(optimizer='adam', loss='mean_squared_error')
```

```
[30]: # Train the model
model.fit(X_train, y_train, epochs=50, batch_size=64, validation_split=0.2)
```

```
Epoch 1/50
12/12 [=====] - 1s 17ms/step - loss: 252.3782 -
val_loss: 259.0193
Epoch 2/50
```

```

12/12 [=====] - 0s 6ms/step - loss: 247.1452 -
val_loss: 254.3651
Epoch 3/50
12/12 [=====] - 0s 5ms/step - loss: 243.3380 -
val_loss: 250.4047
Epoch 4/50
12/12 [=====] - 0s 5ms/step - loss: 238.9238 -
val_loss: 246.8733
Epoch 5/50
12/12 [=====] - 0s 5ms/step - loss: 236.0591 -
val_loss: 243.5087
Epoch 6/50
12/12 [=====] - 0s 5ms/step - loss: 232.5696 -
val_loss: 240.1729
Epoch 7/50
12/12 [=====] - 0s 5ms/step - loss: 228.8369 -
val_loss: 236.8169
Epoch 8/50
12/12 [=====] - 0s 7ms/step - loss: 226.4194 -
val_loss: 233.3858
Epoch 9/50
12/12 [=====] - 0s 5ms/step - loss: 222.1881 -
val_loss: 229.5893
Epoch 10/50
12/12 [=====] - 0s 5ms/step - loss: 218.6411 -
val_loss: 225.3600
Epoch 11/50
12/12 [=====] - 0s 5ms/step - loss: 213.2106 -
val_loss: 220.6633
Epoch 12/50
12/12 [=====] - 0s 4ms/step - loss: 209.2464 -
val_loss: 215.5111
Epoch 13/50
12/12 [=====] - 0s 4ms/step - loss: 204.3389 -
val_loss: 209.9335
Epoch 14/50
12/12 [=====] - 0s 6ms/step - loss: 198.9249 -
val_loss: 203.8828
Epoch 15/50
12/12 [=====] - 0s 5ms/step - loss: 192.0374 -
val_loss: 197.2058
Epoch 16/50
12/12 [=====] - 0s 4ms/step - loss: 185.2454 -
val_loss: 189.8134
Epoch 17/50
12/12 [=====] - 0s 4ms/step - loss: 179.9406 -
val_loss: 181.8738
Epoch 18/50

```

```

12/12 [=====] - 0s 4ms/step - loss: 172.0155 -
val_loss: 173.4478
Epoch 19/50
12/12 [=====] - 0s 5ms/step - loss: 165.0552 -
val_loss: 164.5048
Epoch 20/50
12/12 [=====] - 0s 4ms/step - loss: 153.6075 -
val_loss: 154.9550
Epoch 21/50
12/12 [=====] - 0s 5ms/step - loss: 145.7170 -
val_loss: 144.9520
Epoch 22/50
12/12 [=====] - 0s 6ms/step - loss: 138.1663 -
val_loss: 134.5430
Epoch 23/50
12/12 [=====] - 0s 6ms/step - loss: 126.6693 -
val_loss: 123.8385
Epoch 24/50
12/12 [=====] - 0s 5ms/step - loss: 118.1079 -
val_loss: 113.0589
Epoch 25/50
12/12 [=====] - 0s 5ms/step - loss: 108.6565 -
val_loss: 102.2425
Epoch 26/50
12/12 [=====] - 0s 5ms/step - loss: 97.0588 - val_loss:
91.5823
Epoch 27/50
12/12 [=====] - 0s 5ms/step - loss: 90.7468 - val_loss:
81.2450
Epoch 28/50
12/12 [=====] - 0s 5ms/step - loss: 80.4169 - val_loss:
71.7176
Epoch 29/50
12/12 [=====] - 0s 5ms/step - loss: 73.8794 - val_loss:
62.7222
Epoch 30/50
12/12 [=====] - 0s 5ms/step - loss: 69.9019 - val_loss:
54.5315
Epoch 31/50
12/12 [=====] - 0s 5ms/step - loss: 60.5964 - val_loss:
47.2249
Epoch 32/50
12/12 [=====] - 0s 5ms/step - loss: 53.9876 - val_loss:
40.9681
Epoch 33/50
12/12 [=====] - 0s 5ms/step - loss: 52.7356 - val_loss:
35.6330
Epoch 34/50

```

12/12 [=====] - 0s 6ms/step - loss: 54.1612 - val\_loss:  
31.3898  
Epoch 35/50  
12/12 [=====] - 0s 5ms/step - loss: 48.9563 - val\_loss:  
27.9777  
Epoch 36/50  
12/12 [=====] - 0s 5ms/step - loss: 48.6877 - val\_loss:  
25.1049  
Epoch 37/50  
12/12 [=====] - 0s 5ms/step - loss: 44.6709 - val\_loss:  
23.0617  
Epoch 38/50  
12/12 [=====] - 0s 5ms/step - loss: 44.5108 - val\_loss:  
21.4052  
Epoch 39/50  
12/12 [=====] - 0s 5ms/step - loss: 44.2013 - val\_loss:  
19.7517  
Epoch 40/50  
12/12 [=====] - 0s 5ms/step - loss: 41.0403 - val\_loss:  
18.5326  
Epoch 41/50  
12/12 [=====] - 0s 4ms/step - loss: 43.9875 - val\_loss:  
17.8060  
Epoch 42/50  
12/12 [=====] - 0s 5ms/step - loss: 43.9410 - val\_loss:  
16.9510  
Epoch 43/50  
12/12 [=====] - 0s 5ms/step - loss: 45.4846 - val\_loss:  
16.3857  
Epoch 44/50  
12/12 [=====] - 0s 4ms/step - loss: 40.5152 - val\_loss:  
16.1031  
Epoch 45/50  
12/12 [=====] - 0s 4ms/step - loss: 38.0606 - val\_loss:  
15.6306  
Epoch 46/50  
12/12 [=====] - 0s 4ms/step - loss: 42.5571 - val\_loss:  
15.1468  
Epoch 47/50  
12/12 [=====] - 0s 5ms/step - loss: 41.5191 - val\_loss:  
14.8363  
Epoch 48/50  
12/12 [=====] - 0s 4ms/step - loss: 42.1334 - val\_loss:  
14.5936  
Epoch 49/50  
12/12 [=====] - 0s 4ms/step - loss: 42.1698 - val\_loss:  
14.7882  
Epoch 50/50



```
12/12 [=====] - 0s 5ms/step - loss: 38.6451 - val_loss: 14.9994
```

```
[30]: <keras.callbacks.History at 0x19a8c1f6260>
```

```
[31]: # Make predictions on the test set
y_pred = model.predict(X_test)
```

```
8/8 [=====] - 0s 1ms/step
```

```
[32]: # Calculate the evaluation metrics
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
mse = mean_squared_error(y_test, y_pred)
```

```
[33]: mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
```

```
[34]: # Print the evaluation metrics for the Artificial Neural Network model
print('Artificial Neural Network Model:')
print('RMSE:', rmse)
print('MSE:', mse)
print('MAE:', mae)
print('R-squared:', r2)
```

```
Artificial Neural Network Model:
RMSE: 3.678429963996212
MSE: 13.530847000025176
MAE: 2.9567028927979604
R-squared: 0.4610289809575676
```

#### 0.0.4 Create an XGBoost Regression model

```
[35]: xgb_model = XGBRegressor(objective='reg:squarederror', random_state=42)
```

```
[36]: # Evaluate the model using 5-fold cross-validation
xgb_scores = cross_val_score(xgb_model, X, y, cv=cv,
    ↪scoring='neg_mean_squared_error')
```

```
[37]: # Calculate the evaluation metrics from the scores
xgb_rmse = np.sqrt(-xgb_scores.mean())
xgb_mse = -xgb_scores.mean()
```

```
[38]: xgb_mae = -cross_val_score(xgb_model, X, y, cv=cv,
    ↪scoring='neg_mean_absolute_error').mean()
xgb_r2 = cross_val_score(xgb_model, X, y, cv=cv, scoring='r2').mean()
```

```
[39]: # Print the evaluation metrics for the XGBoost Regression model
print('XGBoost Regression Model:')
print('RMSE:', xgb_rmse)
```

```
print('MSE:', xgb_mse)
print('MAE:', xgb_mae)
print('R-squared:', xgb_r2)
```

XGBoost Regression Model:  
RMSE: 1.6957990140851462  
MSE: 2.875734296172154  
MAE: 1.318473382554611  
R-squared: 0.8921737087350579

```
[40]: import seaborn as sns
```

```
[41]: sns.set_style("whitegrid")
```

```
# Create the lmplo with height=8 and aspect=1.5
sns.lmplot(x='Input1', y='Output', data=data1, scatter_kws={'color': 'red'},
           line_kws={'color': 'blue'}, height=8, aspect=1.5)

# Show the plot
plt.show()
```

